

GROUP NUMBER: 6442

MScFE 652 RISK MANAGEMENT

Group Work Project # 1

Tasks

Step 1 & Step 2

The problem that the thesis attempts to solve is the problem of underlying factors driving crude oil prices. Apart from predicting prices, it is also necessary to understand how crude oil prices react to hypothetical scenarios like geopolitical tensions or sudden changes in supply. Therefore, the thesis attempts to solve this problem by performing stress tests and evaluating the robustness of the model under extreme conditions.

Bayesian networks are suitable for solving the problem of underlying factors driving crude oil prices because it can combine a lot of variables to find the dependencies between them and successfully extrapolate the data thanks to the learned knowledge. This method applies probabilistic reasoning and can be dynamically adapted towards an ever changing world. Bayesian networks models casual relationships between variables which allows for effective stress testing by simulating the impact of different hypothetical situations on crude oil prices.

The main advantage is that it lets us introduce prior knowledge that could be either expert-based or extracted from other models and put it into the decision process. As compared to traditional time-series models that might overlook the complex interactions like the underlying factors driving crude oil prices, bayesian networks provide more accurate and reliable forecasts.

Step 3 Student A

```
In [1]: !pip install fredapi
        !pip install hmms
```

Collecting fredapi

Downloading fredapi-0.5.2-py3-none-any.whl (11 kB)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from fredapi) (2.0.3)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (2023.4)

Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (2024.1)

Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (1.25.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->fredapi) (1.16.0)

Installing collected packages: fredapi

Successfully installed fredapi-0.5.2

Collecting hmms

Downloading hmms-0.2.3.tar.gz (524 kB)

524.8/524.8 kB 3.6 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

Requirement already satisfied: Cython in /usr/local/lib/python3.10/dist-packages (from hmms) (3.0.10)

Requirement already satisfied: NumPy in /usr/local/lib/python3.10/dist-packages (from hmms) (1.25.2)

Requirement already satisfied: ipython in /usr/local/lib/python3.10/dist-packages (from hmms) (7.34.0)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from hmms) (3.7.1)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from hmms) (2.0.3)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from hmms) (1.11.4)

Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (67.7.2)

Collecting jedi>=0.16 (from ipython->hmms)

Downloading jedi-0.19.1-py2.py3-none-any.whl (1.6 MB)

1.6/1.6 MB 24.7 MB/s eta 0:00:00

Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (4.4.2)

Requirement already satisfied: pickleshare in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (0.7.5)

Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (5.7.1)

Requirement already satisfied: prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (3.0.47)

Requirement already satisfied: pygments in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (2.16.1)

Requirement already satisfied: backcall in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (0.2.0)

Requirement already satisfied: matplotlib-inline in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (0.1.7)

Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.10/dist-packages (from ipython->hmms) (4.9.0)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (1.2.1)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages

```
(from matplotlib->hmms) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->hmms) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hmms) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hmms) (2024.1)
Requirement already satisfied: parso<0.9.0,>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from jedi>=0.16->ipython->hmms) (0.8.4)
Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.10/dist-packages (from pexpect>4.3->ipython->hmms) (0.7.0)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0->ipython->hmms) (0.2.13)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib->hmms) (1.16.0)
Building wheels for collected packages: hmms
  Building wheel for hmms (setup.py) ... done
  Created wheel for hmms: filename=hmms-0.2.3-cp310-cp310-linux_x86_64.whl size=2152616 sha256=06160f72bea8bab201c598c6c29a64cb06b43c03cb3fe901a7d71824cc25306b
  Stored in directory: /root/.cache/pip/wheels/aa/6f/a4/1dbae244341f24881dce9465aa533729d2ae870cff3866070f
Successfully built hmms
Installing collected packages: jedi, hmms
Successfully installed hmms-0.2.3 jedi-0.19.1
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
import statsmodels.api as sm
from scipy import stats
from statsmodels.tsa.stattools import kpss, adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.stats.diagnostic import acorr_ljungbox
```

```
In [3]: from fredapi import Fred
# FRED API key
fred_key = "f9c22fca078ece81a7a2ac6fba29b8a9";
# Initiates a session with the FRED datacenter to retrieve datasets
fred = Fred(api_key=fred_key);
# Retrieve data from FRED API
fred_data = pd.DataFrame(fred.get_series('WTISPLC'), columns=['WTISPLC'])
```

EDA

```
In [4]: datasets_fred = [
    'WTISPLC',    # Spot Crude Oil Price: West Texas Intermediate (WTI) (WTISPLC)
    'CPIENGSL',   # Consumer Price Index for All Urban Consumers: Energy in U.S. City Average
    'CAPG211S',   # Industrial Capacity: Mining: Oil and Gas Extraction (NAICS = 211)
    'CAPUTLG211S', # Capacity Utilization: Mining: Oil and Gas Extraction (NAICS = 211)
    'IPG211S',    # Industrial Production Index: Mining: Oil and Gas Extraction (NAICS = 211)
    'INDPRO',     # Industrial Production: Total Index
    'IPN213111N', # Industrial Production: Mining: Drilling Oil and Gas Wells
    'PCU211211',  # Producer Price Index: Mining: Oil and Gas Extraction (NAICS = 211)
    ];

data_frames = []; # List of dataframes to be concatenated

# Adding FRED datasets
for series_id in datasets_fred:
    # Get series from FRED
    df = pd.DataFrame(fred.get_series(series_id), columns=[series_id]);
    data_frames.append(df);

data_merge = pd.concat(data_frames, axis=1)
#data_merge = data_merge[data_merge.index > '2000-01-01']
data_merge.head()
```

Out[4]:

	WTISPLC	CPIENGSL	CAPG211S	CAPUTLG211S	IPG211S	INDPRO	IPN213111N	PCI
--	---------	----------	----------	-------------	---------	--------	------------	-----

1919-01-01	NaN	NaN	NaN	NaN	NaN	4.8654	NaN	
------------	-----	-----	-----	-----	-----	--------	-----	--

1919-02-01	NaN	NaN	NaN	NaN	NaN	4.6504	NaN	
------------	-----	-----	-----	-----	-----	--------	-----	--

1919-03-01	NaN	NaN	NaN	NaN	NaN	4.5160	NaN	
------------	-----	-----	-----	-----	-----	--------	-----	--

1919-04-01	NaN	NaN	NaN	NaN	NaN	4.5966	NaN	
------------	-----	-----	-----	-----	-----	--------	-----	--

1919-05-01	NaN	NaN	NaN	NaN	NaN	4.6235	NaN	
------------	-----	-----	-----	-----	-----	--------	-----	--



```
In [5]: data_merge.isnull().sum()
```

```
Out[5]: WTISPLC      324
        CPIENGSL     456
        CAPG211S     637
        CAPUTLG211S  637
        IPG211S      637
        INDPRO        1
        IPN213111N   637
        PCU211211    804
        dtype: int64
```

```
In [6]: data_merge.describe().T
```

```
Out[6]:
```

	count	mean	std	min	25%	50%	75%	m
WTISPLC	942.0	27.816605	29.284847	1.1700	3.0000	18.3125	39.34000	133.93
CPIENGSL	810.0	114.829889	80.335992	21.3000	31.7500	101.5000	192.18225	331.73
CAPG211S	629.0	81.421457	21.937600	61.4822	66.6233	76.3514	81.75120	149.28
CAPUTLG211S	629.0	93.202929	3.239208	78.6360	91.5201	93.1402	94.95650	101.54
IPG211S	629.0	76.173594	22.164265	48.8141	62.4444	68.5466	78.22200	144.42
INDPRO	1265.0	45.891390	34.813005	3.6827	13.7629	39.1057	84.15640	104.10
IPN213111N	629.0	130.404341	53.132521	47.9947	94.0882	113.9683	156.98800	334.62
PCU211211	462.0	154.771381	83.664228	54.6000	77.3750	138.1500	221.97500	490.40

Step 3 Student C

Identified that the financial data on securities and markets that affect the price of oil are:

1. Canadian dollar to US dollar exchange rate
2. CBOE crude oil futures
3. WTI crude oil futures
4. Brent crude oil futures
5. S&P500

Importing, structuring and graphing the financial data on securities and markets that affect the price of oil.

```
In [7]: from fredapi import Fred
        # FRED API key
        fred_key = "f9c22fca078ece81a7a2ac6fba29b8a9";
        # Initiates a session with the FRED datacenter to recieve datasets
        fred = Fred(api_key=fred_key);
        # Retrieve data from FRED API
        fred_data = pd.DataFrame(fred.get_series('DEXCAUS'), columns=['DEXCAUS'])
```

```

In [8]: financial_datasets_fred = [
        'DEXCAUS', #Canadian dollar to US dollar exchange rate
        'VIXCLS', #CBOE Volatility Index
        'DCOILWTICO', #WTI Crude oil futures
        'DCOILBRETEU', #Brent crude oil futures
        'SP500', #S&P500 Index
        ];

        finan_data_frames = []; # List of dataframes to be concatenated

        # Adding FRED datasets
        for series_id in financial_datasets_fred:
            # Get series from FRED
            df = pd.DataFrame(fred.get_series(series_id), columns=[series_id]);
            finan_data_frames.append(df);

        fin_data = pd.concat(finan_data_frames, axis=1)
        #fin_data = data_merge[fin_data.index > '2000-01-01']
        fin_data.head()

```

```

Out[8]:

```

	DEXCAUS	VIXCLS	DCOILWTICO	DCOILBRETEU	SP500
1971-01-04	1.0109	NaN	NaN	NaN	NaN
1971-01-05	1.0102	NaN	NaN	NaN	NaN
1971-01-06	1.0106	NaN	NaN	NaN	NaN
1971-01-07	1.0148	NaN	NaN	NaN	NaN
1971-01-08	1.0154	NaN	NaN	NaN	NaN

```

In [9]: fin_data.isnull().sum()

```

```

Out[9]: DEXCAUS          539
        VIXCLS          5253
        DCOILWTICO      4264
        DCOILBRETEU     4541
        SP500          11448
        dtype: int64

```

```

In [51]: fin_data.describe().T

```

Out[51]:

	count	mean	std	min	25%	50%	75%
DEXCAUS	13425.0	1.229735	0.160415	0.9168	1.1035	1.2353	1.34
VIXCLS	8711.0	19.491915	7.879984	9.1400	13.7950	17.6300	22.87
DCOILWTICO	9700.0	47.291728	29.700010	-36.9800	20.1500	39.2800	70.94
DCOILBRETEU	9423.0	49.936190	32.927056	9.1000	19.2800	42.7200	74.23
SP500	2516.0	3177.629996	979.254402	1829.0800	2268.9750	2890.6550	4090.42

Step 4

Dictionary of Macroeconomic data

```
In [11]: start_date = data_merge.index.min()
end_date = data_merge.index.max()

macro_data_info = {
    'WTISPLC': {
        'description': 'Oil Price',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'CPIENGSL': {
        'description': 'Consumer Price Index for All Urban Consumers: Energy in U.S.',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'CAPG211S': {
        'description': 'Industrial Capacity: Mining: Oil and Gas Extraction',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'CAPUTLG211S': {
        'description': 'Capacity Utilization: Mining: Oil and Gas Extraction',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'IPG211S': {
        'description': 'Industrial Production Index: Mining: Oil and Gas Extraction',
        'frequency': 'Monthly',
```

```

        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'INDPRO': {
        'description': 'Industrial Production: Total Index',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'IPN213111N': {
        'description': 'Industrial Production: Mining: Drilling Oil and Gas Wells',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'PCU211211': {
        'description': 'Producer Price Index: Mining: Oil and Gas Extraction',
        'frequency': 'Monthly',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    }
}

macro_data_info_list = [
    ['WTISPLC', 'Spot Crude Oil Price: West Texas Intermediate (WTI)', 'Monthly', 'FRED', start_date, end_date],
    ['CPIENGSL', 'Consumer Price Index for All Urban Consumers: Energy in U.S. City Average', 'Monthly', 'FRED', start_date, end_date],
    ['CAPG211S', 'Industrial Capacity: Mining: Oil and Gas Extraction', 'Monthly', 'FRED', start_date, end_date],
    ['CAPUTLG211S', 'Capacity Utilization: Mining: Oil and Gas Extraction', 'Monthly', 'FRED', start_date, end_date],
    ['IPG211S', 'Industrial Production Index: Mining: Oil and Gas Extraction', 'Monthly', 'FRED', start_date, end_date],
    ['INDPRO', 'Industrial Production: Total Index', 'Monthly', 'FRED', start_date, end_date],
    ['IPN213111N', 'Industrial Production: Mining: Drilling Oil and Gas Wells', 'Monthly', 'FRED', start_date, end_date],
    ['PCU211211', 'Producer Price Index: Mining: Oil and Gas Extraction', 'Monthly', 'FRED', start_date, end_date]
]

macro_data_info = pd.DataFrame(macro_data_info_list, columns=['Ticker',
                                                             'Description',
                                                             'Frequency',
                                                             'Source',
                                                             'Start Date',
                                                             'End Date'])

macro_data_info

```


Out[11]:

	Ticker	Description	Frequency	Source	Start Date	End Date
0	WTISPLC	Spot Crude Oil Price: West Texas Intermediate ...	Monthly	FRED	1919-01-01	2024-06-01
1	CPIENGSL	Consumer Price Index for All Urban Consumers: ...	Monthly	FRED	1919-01-01	2024-06-01
2	CAPG211S	Industrial Capacity: Mining: Oil and Gas Extra...	Monthly	FRED	1919-01-01	2024-06-01
3	CAPUTLG211S	Capacity Utilization: Mining: Oil and Gas Extr...	Monthly	FRED	1919-01-01	2024-06-01
4	IPG211S	Industrial Production Index: Mining: Oil and G...	Monthly	FRED	1919-01-01	2024-06-01
5	INDPRO	Industrial Production: Total Index	Monthly	FRED	1919-01-01	2024-06-01
6	IPN213111N	Industrial Production: Total Index	Monthly	FRED	1919-01-01	2024-06-01
7	PCU211211	Producer Price Index: Mining: Oil and Gas Extr...	Monthly	FRED	1919-01-01	2024-06-01

Dictionary of the financial data:

```
In [12]: start_date = fin_data.index.min()
end_date = fin_data.index.max()

print("Start Date:", start_date)
print("End Date:", end_date)
```

Start Date: 1971-01-04 00:00:00
End Date: 2024-07-11 00:00:00

```
In [13]: data_info = {
    'DEXCAUS': {
        'description': 'Canadian dollar to US dollar exchange rate',
        'frequency': 'Daily',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'VIXCLS': {
        'description': 'CBOE Volatility Index',
        'frequency': 'Daily',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'DCOILWTICO': {
        'description': 'WTI Crude oil futures',
        'frequency': 'Daily',
```

```

        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'DCOILBRETEU': {
        'description': 'Brent crude oil futures',
        'frequency': 'Daily',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'SP500-45': {
        'description': 'S&P 500 Energy sector index',
        'frequency': 'Daily',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    },
    'SP500': {
        'description': 'S&P500 Index',
        'frequency': 'Daily',
        'source': 'FRED',
        'start_date': start_date,
        'end_date': end_date
    }
}

# Print the dictionary
for key, value in data_info.items():
    print(f"{key}: {value}")

```

```

DEXCAUS: {'description': 'Canadian dollar to US dollar exchange rate', 'frequency': 'Daily', 'source': 'FRED', 'start_date': Timestamp('1971-01-04 00:00:00'), 'end_date': Timestamp('2024-07-11 00:00:00')}
VIXCLS: {'description': 'CBOE Volatility Index', 'frequency': 'Daily', 'source': 'FRED', 'start_date': Timestamp('1971-01-04 00:00:00'), 'end_date': Timestamp('2024-07-11 00:00:00')}
DCOILWTICO: {'description': 'WTI Crude oil futures', 'frequency': 'Daily', 'source': 'FRED', 'start_date': Timestamp('1971-01-04 00:00:00'), 'end_date': Timestamp('2024-07-11 00:00:00')}
DCOILBRETEU: {'description': 'Brent crude oil futures', 'frequency': 'Daily', 'source': 'FRED', 'start_date': Timestamp('1971-01-04 00:00:00'), 'end_date': Timestamp('2024-07-11 00:00:00')}
SP500-45: {'description': 'S&P 500 Energy sector index', 'frequency': 'Daily', 'source': 'FRED', 'start_date': Timestamp('1971-01-04 00:00:00'), 'end_date': Timestamp('2024-07-11 00:00:00')}
SP500: {'description': 'S&P500 Index', 'frequency': 'Daily', 'source': 'FRED', 'start_date': Timestamp('1971-01-04 00:00:00'), 'end_date': Timestamp('2024-07-11 00:00:00')}

```

Table showing the financial data

```

In [14]: data_info_list = [
    ['DEXCAUS', 'Canadian dollar to US dollar exchange rate', 'Daily', 'FRED', start_date, end_date],
    ['VIXCLS', 'CBOE Volatility Index', 'Daily', 'FRED', start_date, end_date],
    ['DCOILWTICO', 'WTI Crude oil futures', 'Daily', 'FRED', start_date, end_date],
    ['DCOILBRETEU', 'Brent crude oil futures', 'Daily', 'FRED', start_date, end_date]
]

```

```
['SP500-45', 'S&P 500 Energy sector index', 'Daily', 'FRED', start_date, end_date]
['SP500', 'S&P500 Index', 'Daily', 'FRED', start_date, end_date]
]

data_info_df = pd.DataFrame(data_info_list, columns=['Ticker', 'Description', 'Frequency', 'Source', 'Start Date', 'End Date'])

data_info_df
```

Out[14]:

	Ticker	Description	Frequency	Source	Start Date	End Date
0	DEXCAUS	Canadian dollar to US dollar exchange rate	Daily	FRED	1971-01-04	2024-07-11
1	VIXCLS	CBOE Volatility Index	Daily	FRED	1971-01-04	2024-07-11
2	DCOILWTICO	WTI Crude oil futures	Daily	FRED	1971-01-04	2024-07-11
3	DCOILBRENTU	Brent crude oil futures	Daily	FRED	1971-01-04	2024-07-11
4	SP500-45	S&P 500 Energy sector index	Daily	FRED	1971-01-04	2024-07-11
5	SP500	S&P500 Index	Daily	FRED	1971-01-04	2024-07-11

Step 5

Removing Outliers

Macroeconomic Data

```
In [15]: def out_std(df, column):
    global lower, upper
    # calculate the mean and standard deviation of the data frame
    data_mean, data_std = df[column].mean(), df[column].std()
    # calculate the cutoff value
    cut_off = data_std * 3
    # calculate the lower and upper bound value
    lower, upper = data_mean - cut_off, data_mean + cut_off
    print('The lower bound value is', lower)
    print('The upper bound value is', upper)
    # Calculate the number of records below and above lower and above bound value respectively
    df1 = df[df[column] > upper]
    df2 = df[df[column] < lower]
    return print('Total number of outliers are', df1.shape[0] + df2.shape[0])
```

```

for col in data_merge.columns:
    out_std(data_merge, col)

plt.figure(figsize = (10,6))
sns.distplot(data_merge[col], kde=False)
plt.axvspan(xmin = lower,xmax= data_merge[col].min(),alpha=0.2, color='red')
plt.axvspan(xmin = upper,xmax= data_merge[col].max(),alpha=0.2, color='red')

```

The lower bound value is -60.037936242848204

The upper bound value is 115.67114643393101

Total number of outliers are 4

The lower bound value is -126.1780877872839

The upper bound value is 355.83786556506163

Total number of outliers are 0

The lower bound value is 15.608656858545089

The upper bound value is 147.23425697293345

Total number of outliers are 6

<ipython-input-15-238fc34c1aa0>:20: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(data_merge[col], kde=False)
```

The lower bound value is 83.48530555873768

The upper bound value is 102.92055199293165

Total number of outliers are 5

The lower bound value is 9.680800381005028

The upper bound value is 142.6663878542891

Total number of outliers are 9

The lower bound value is -58.54762458917536

The upper bound value is 150.33040419391844

Total number of outliers are 0

The lower bound value is -28.99322264970334



The upper bound value is 289.8019050026445



Total number of outliers are 12



The lower bound value is -96.22130158117326



The upper bound value is 405.76406348593514

Total number of outliers are 3









```
In [16]: df_new = pd.DataFrame()

for col in data_merge.columns:
    #Data Frame without outliers
    df_new = data_merge[(data_merge[col] > lower) & (data_merge[col] < upper) ]

for col in df_new.columns:
    out_std(df_new, col)
```



```
plt.figure(figsize = (10,6))
sns.distplot(df_new[col], kde=False)
plt.axvspan(xmin = lower,xmax= df_new[col].min(),alpha=0.2, color='red')
plt.axvspan(xmin = upper,xmax= df_new[col].max(),alpha=0.2, color='red')
```

The lower bound value is -40.082532829843636

The upper bound value is 133.40170930043186

Total number of outliers are 0

The lower bound value is -23.57509455720819

The upper bound value is 357.58868061385306

Total number of outliers are 0

The lower bound value is 6.15368154560106

The upper bound value is 159.17391540429003

Total number of outliers are 0

The lower bound value is 82.76136514142729

The upper bound value is 104.20412113308254

Total number of outliers are 5

The lower bound value is 0.5550927639357894

The upper bound value is 154.7436218330141

Total number of outliers are 0

<ipython-input-16-d512fc6a9f76>:11: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_new[col], kde=False)
```

The lower bound value is 39.572396006374206

The upper bound value is 133.53064015920313

Total number of outliers are 0

The lower bound value is 23.142984182915242



The upper bound value is 196.66759620924165



Total number of outliers are 0



The lower bound value is -87.6888973772983



The upper bound value is 393.2304137171676

Total number of outliers are 0









Financial Data

```
In [17]: def out_std(df, column):
          global lower, upper
          # calculate the mean and standard deviation of the data frame
          data_mean, data_std = df[column].mean(), df[column].std()
          # calculate the cutoff value
          cut_off = data_std * 3
          # calculate the lower and upper bound value
```

```

lower, upper = data_mean - cut_off, data_mean + cut_off
print('The lower bound value is', lower)
print('The upper bound value is', upper)
# Calculate the number of records below and above lower and above bound value res
df1 = df[df[column] > upper]
df2 = df[df[column] < lower]
return print('Total number of outliers are', df1.shape[0]+ df2.shape[0])

for col in fin_data.columns:
    out_std(fin_data, col)

plt.figure(figsize = (10,6))
sns.distplot(fin_data[col], kde=False)
plt.axvspan(xmin = lower,xmax= fin_data[col].min(),alpha=0.2, color='red')
plt.axvspan(xmin = upper,xmax= fin_data[col].max(),alpha=0.2, color='red')

```

The lower bound value is 0.7484901853824507
 The upper bound value is 1.7109794459024652
 Total number of outliers are 0
 The lower bound value is -4.148036317968636
 The upper bound value is 43.131865958652824
 Total number of outliers are 142
 The lower bound value is -41.80830253900812
 The upper bound value is 136.3917582091112
 Total number of outliers are 16

<ipython-input-17-a4129b99bc3f>:20: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>



```


sns.distplot(fin_data[col], kde=False)

```

The lower bound value is -48.84497754215096
 The upper bound value is 148.71735788811296
 Total number of outliers are 0
 The lower bound value is 239.8667915078745
 The upper bound value is 6115.393200543
 Total number of outliers are 0







Handling missing values and Step 6

```
In [18]: clean_filtered_macrodata = df_new[df_new.index > '2000-01-01']
for col in clean_filtered_macrodata.columns:
    #Data Frame without outliers
    clean_filtered_macrodata = clean_filtered_macrodata[(clean_filtered_macrodata[col]
clean_filtered_macrodata.dropna())
```

Out[18]: **WTISPLC CPIENGL CAPG211S CAPUTLG211S IPG211S INDPRO IPN213111N PCU2112**



```
In [19]: clean_filtered_fin_data = fin_data[fin_data.index > '2000-01-01']

for col in clean_filtered_fin_data.columns:
    out_std(clean_filtered_fin_data, col)
    clean_filtered_fin_data = clean_filtered_fin_data[(clean_filtered_fin_data[col] >

clean_filtered_fin_data
```

The lower bound value is 0.7383116741499283
 The upper bound value is 1.754145663956157
 Total number of outliers are 0
 The lower bound value is -5.619809271550544
 The upper bound value is 45.44014309580189
 Total number of outliers are 102
 The lower bound value is -13.297905292740182
 The upper bound value is 140.99957836003023
 Total number of outliers are 8
 The lower bound value is -20.30716118329657
 The upper bound value is 153.93805280642147
 Total number of outliers are 0
 The lower bound value is 241.3547172663093
 The upper bound value is 6106.041783351229
 Total number of outliers are 0

Out[19]:

	DEXCAUS	VIXCLS	DCOILWTICO	DCOILBRETEU	SP500
2014-07-14	1.0721	11.82	101.73	104.73	1977.10
2014-07-15	1.0766	11.96	100.56	104.73	1973.28
2014-07-16	1.0749	11.00	101.88	105.41	1981.57
2014-07-17	1.0753	14.54	103.84	106.04	1958.12
2014-07-18	1.0737	12.06	103.83	106.03	1978.22
...
2024-06-28	1.3684	12.44	82.83	87.26	5460.48
2024-07-01	1.3742	12.22	84.70	86.57	5475.09
2024-07-02	1.3691	12.03	84.09	88.28	5509.01
2024-07-03	1.3624	12.09	85.19	88.25	5537.02
2024-07-05	1.3631	12.48	84.44	88.66	5567.19

2429 rows × 5 columns

In [20]: `clean_filtered_fin_data.isnull().sum()`

Out[20]: DEXCAUS 0
 VIXCLS 0
 DCOILWTICO 0
 DCOILBRETEU 0
 SP500 0
 dtype: int64

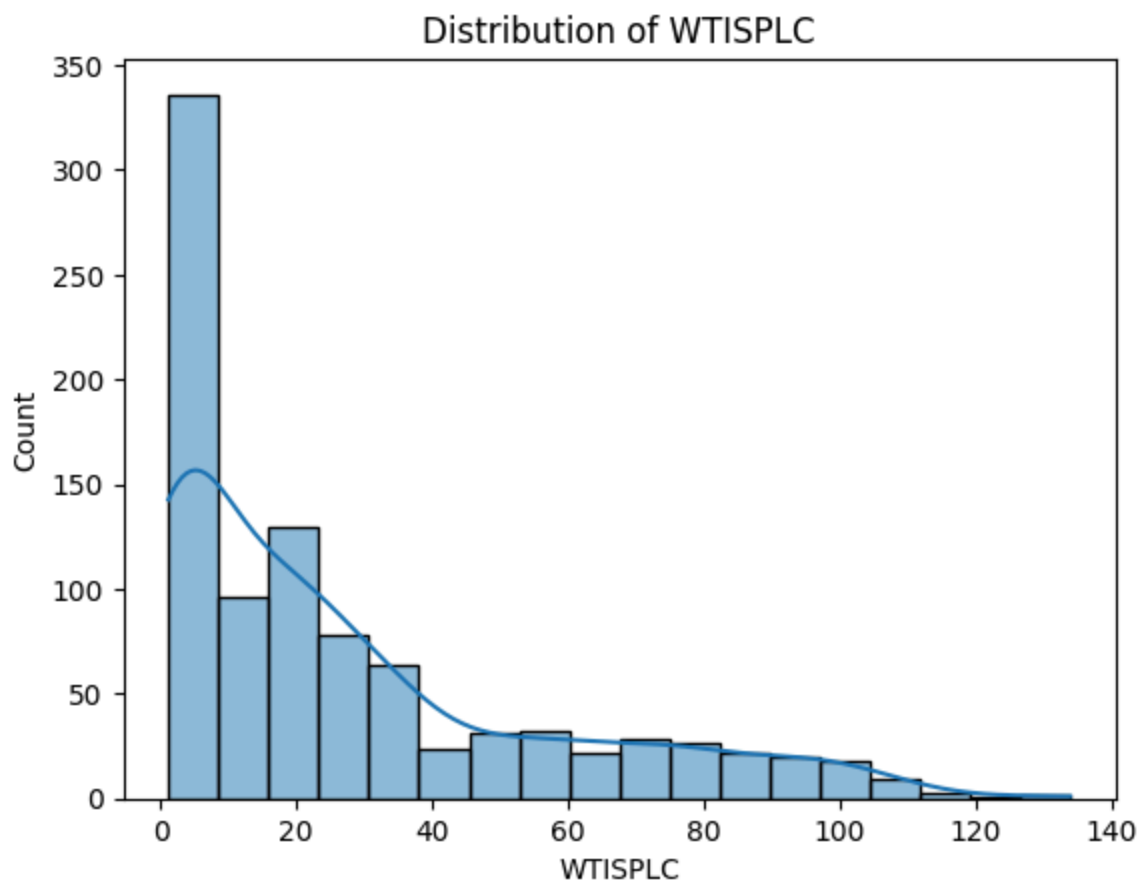
We are filtering all the data before 01.01.2000 and from now on we are going to work with only the data after 2000 because we think that this timeframe is big enough and also this is the most recent 24 years that do have the most and recent important impact over history and the corresponding price action. In order to to get rid of the outliers we are applying the 3 standard deviation outlier detection method and after this we are dropping the data point

that are outside of the 3 standard deviation metric. After this we check for null data and we gladly find out that there is not so no further imputation is needed.


Step 7

Distributional plots of Macroeconomic Data


```
In [21]: # Distribution of each variable
for col in data_merge.columns:
    plt.figure()
    sns.histplot(data_merge[col], kde=True)
    plt.title(f"Distribution of {col}")
    plt.show()
```




Distribution of CPIENGSL




Distribution of CAPG211S


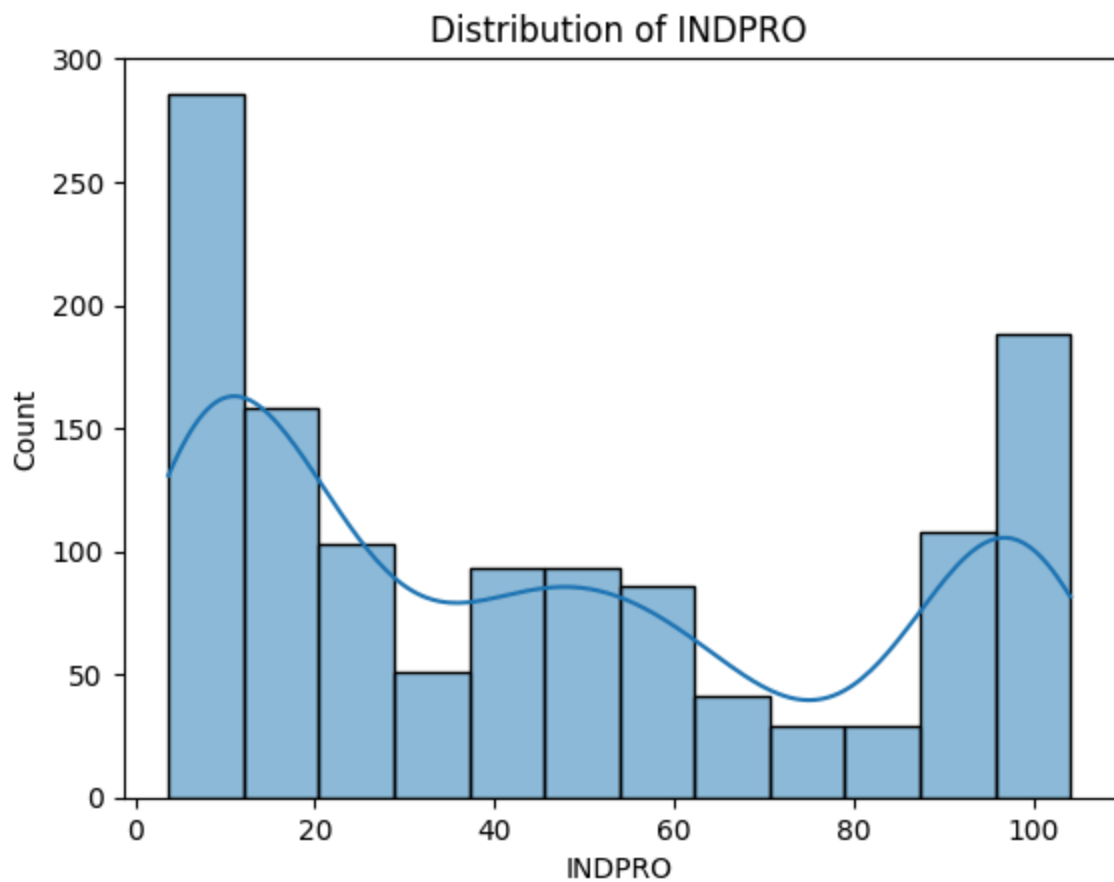



Distribution of CAPUTLG211S



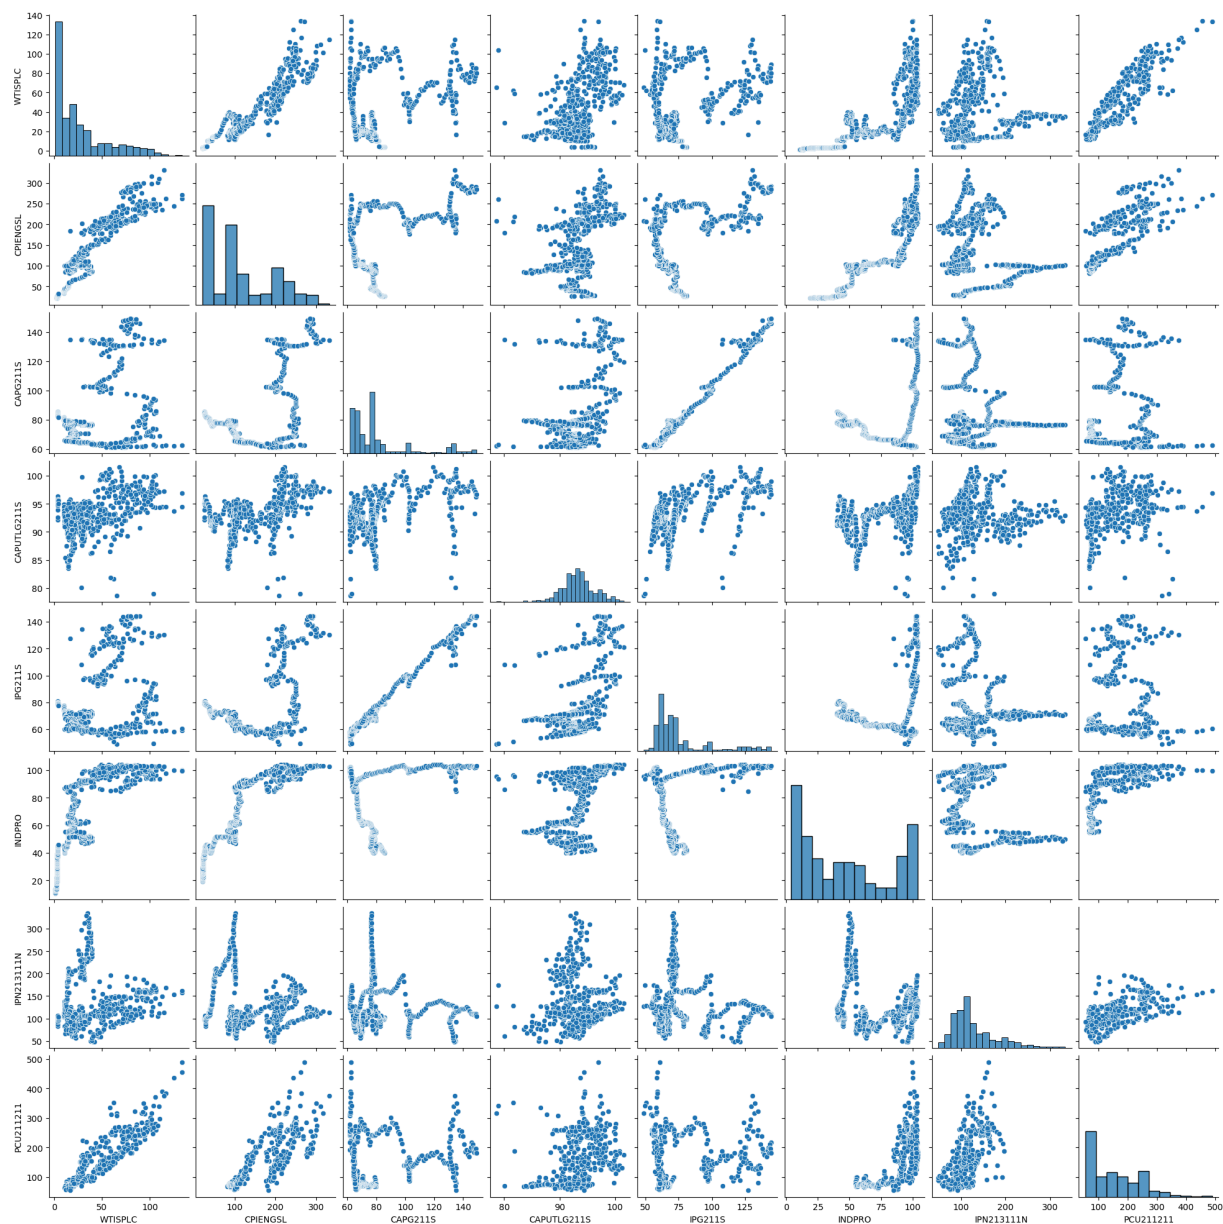
Distribution of IPG211S








```
In [22]: # Scatter plot matrix  
sns.pairplot(data_merge)  
plt.show()
```




```
In [23]: # Correlation matrix
corr_matrix = data_merge.corr()
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.show()
```

Time Series plot of Macroeconomic Data


```
In [24]: # Time series plot
data_merge.plot(figsize=(10, 6))
plt.xlabel("Date")
plt.ylabel("Value")
plt.title("Time Series Plot of Data")
plt.show()
```




```
In [25]: 'WTISPLC',    # Spot Crude Oil Price: West Texas Intermediate (WTI) (WTISPLC)
          'CPIENGSL',  # Consumer Price Index for ALL Urban Consumers: Energy in U.S. City Average
          'CAPG211S',  # Industrial Capacity: Mining: Oil and Gas Extraction (NAICS = 211)
          'CAPUTLG211S', # Capacity Utilization: Mining: Oil and Gas Extraction (NAICS = 211)
          'IPG211S',   # Industrial Production Index: Mining: Oil and Gas Extraction (NAICS = 211)
          'INDPRO',    # Industrial Production: Total Index
          'IPN213111N', # Industrial Production: Mining: Drilling Oil and Gas Wells
          'PCU211211', # Producer Price Index: Mining: Oil and Gas Extraction (NAICS = 211)
```

```
Out[25]: ('PCU211211',)
```


```
In [26]: # Time series plot
data_merge['WTISPLC'].dropna().plot(figsize=(10, 6)) # Spot Crude Oil Price: West Texas Intermediate
plt.xlabel("Date")
plt.ylabel("$")
plt.title("Oil Price Dynamics")
plt.show()
```




```
In [27]: # Time series plot
data_merge['CPIENGSL'].dropna().plot(figsize=(10, 6)) # Consumer Price Index for ALL
plt.xlabel("Date")
plt.ylabel("Value")
plt.title("Consumer Price Index for All Urban Consumers: Energy in U.S. City Average")
plt.show()
```




```
In [28]: # Time series plot
data_merge['CAPG211S'].dropna().plot(figsize=(10, 6)) # Industrial Capacity: Mining:
plt.xlabel("Date")
plt.ylabel("Quantity")
plt.title("Industrial Capacity: Mining: Oil and Gas Extraction ")
plt.show()
```




```
In [29]: # Time series plot
data_merge['INDPRO'].dropna().plot(figsize=(10, 6)) # Industrial Production: Total
plt.xlabel("Date")
plt.ylabel("Quantity")
plt.title(" Industrial Production: Total Index")
plt.show()
```



```
In [30]: # Time series plot
data_merge['IPN213111N'].dropna().plot(figsize=(10, 6)) # Industrial Production: Mining: Drilling Oil and Gas Wells
plt.xlabel("Date")
plt.ylabel("Quantity")
plt.title(" Industrial Production: Mining: Drilling Oil and Gas Wells ")
plt.show()
```




```
In [31]: # Time series plot
data_merge['PCU211211'].dropna().plot(figsize=(10, 6)) # Producer Price Index: Mining
plt.xlabel("Date")
plt.ylabel("Quantity")
plt.title("Producer Price Index: Mining: Oil and Gas Extraction")
plt.show()
```




Distributional Plot of Financial Data



```
In [32]: # Distribution of each variable
for col in fin_data.columns:
    plt.figure()
    sns.histplot(fin_data[col], kde=True)
    plt.title(f"Distribution of {col}")
    plt.show()
```


Distribution of DEXCAUS



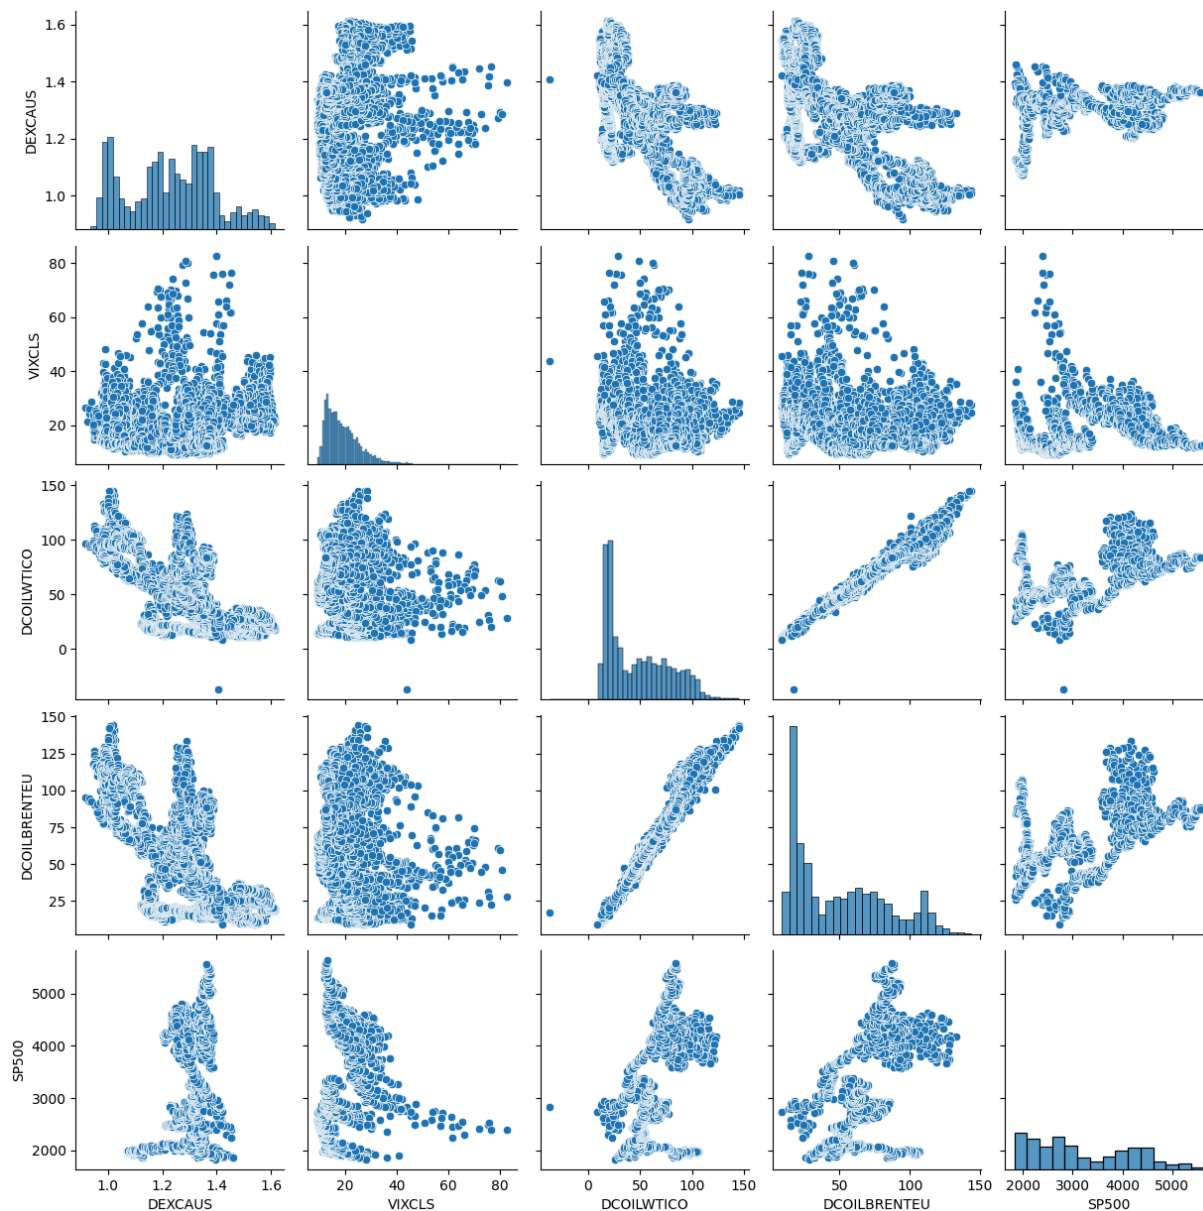
Distribution of VIXCLS








```
In [33]: # Scatter plot matrix  
sns.pairplot(fin_data)  
plt.show()
```




```
In [34]: # Correlation matrix
corr_matrix = fin_data.corr()
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.show()
```




Time Series Plot of Financial Data

```
In [35]: # Time series plot
fin_data.plot(figsize=(10, 6))
plt.xlabel("Date")
plt.ylabel("Value")
plt.title("Time Series Plot of Data")
plt.show()
```




```
In [36]: # Time series plot
fin_data['DEXCAUS'].dropna().plot(figsize=(10, 6)) # Canadian dollar to US dollar ex
plt.xlabel("Date")
plt.ylabel("$")
plt.title("Canadian dollar to US dollar exchange rate")
plt.show()
```




```
In [37]: # Time series plot
fin_data['VIXCLS'].dropna().plot(figsize=(10, 6)) # CBOE Volatility Index
plt.xlabel("Date")
```


```
plt.ylabel("Value")  
plt.title("CBOE Volatility Index")  
plt.show()
```




```
In [38]: # Time series plot  
fin_data['DCOILWTICO'].dropna().plot(figsize=(10, 6)) # WTI Crude oil futures  
plt.xlabel("Date")  
plt.ylabel("Value")  
plt.title("WTI Crude oil futures")  
plt.show()
```



```
In [39]: # Time series plot
fin_data['DCOILBRETEU'].dropna().plot(figsize=(10, 6)) # Brent crude oil futures
plt.xlabel("Date")
plt.ylabel("Value")
plt.title("Brent crude oil futures")
plt.show()
```




```
In [40]: # Time series plot
fin_data['SP500'].dropna().plot(figsize=(10, 6)) # S&P500 Index
plt.xlabel("Date")
plt.ylabel("Value")
plt.title("S&P500 Index")
plt.show()
```




Step 8

```
In [63]: # Time series plot
df_new['WTISPLC'].plot(figsize=(10, 6)) # Spot Crude Oil Price: West Texas Intermedi
plt.xlabel("Date")
plt.ylabel("$")
plt.title("Oil Price Dinamics")
plt.show()
```




We see that the for the period 2000 - 2024 Oil price is locked in a very big range from 20*till* 140 with big spikes and steep declines as well. Usually the big spikes are at periods of booming economy and afterwards followed by crisis that are known with low consumption of oil so it is quite normal for prices to normalize and fall down. Such booming economy periods are 2002 - 2008, 2010 - 2012, 2020 - 2023 and they are followoed crisis (2008 - 2009 The Gread Financial Crisis, 2015-2016, 2019-2020 Covid19) and step decline in demand for oil.

```
In [64]: df_new['WTISPLC'].pct_change().plot(figsize=(10, 6))
plt.xlabel("Date")
plt.ylabel("Percent Change")
plt.title("Oil Price Percent Change")
plt.show()
```




We can say that the plot of Oil prices returns for the period 2000 - 2024 look quite normal (but we have to check it with some statistical methods) with the biggest spike around Covid19. Lets look at the bigger picture: all the data

```
In [43]: data_merge['WTISPLC'].dropna().pct_change().plot(figsize=(10, 6))
plt.xlabel("Date")
plt.ylabel("Percent Change")
plt.title("Oil Price Percent Change")
plt.show()
```

The entire Oil returns data show a more ugly picture with a lot more spikes. Lets see the ACF and Pacf plots for more insights

```
In [44]: # Create the plots
fig, axes = plt.subplots(2, figsize=(12,6))
plot_acf(data_merge['WTISPLC'].dropna(), lags=100, ax=axes[0])
plot_pacf(data_merge['WTISPLC'].dropna(), lags=100, ax=axes[1])
plt.show()
```



From the ACF plot we can see that there is trend in the Oil price, which is quite normal as everything with time goes up thanks to the inflation (a totally normal process). The interesting

thing is the PACF where we can see that there is a seasonality pattern where we can see that at the beginning of the year there is a downward spike (during the northern globe is winter and since there is most of the population located and the fact that it is winter suggests that there is also a lower demand for Oil) and also during the summer of the year (I am talking about the northern globe seasons) then an upward spike in the Oil consumption because it is summer time and a lot of traveling is happening at this time of the year

```
In [45]: # Perform ADF test
print('Results of ADF Test:')
# Replace infinite or nan values with finite numbers
data_merge['WTISPLC'] = data_merge['WTISPLC'].replace([np.inf, -np.inf], np.nan).fill
# Calculate percentage change on the shifted series to avoid introducing NaNs
dfctest = adfuller(data_merge['WTISPLC'].dropna().shift().pct_change().dropna())
dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'P-value', '#Lags Used', 'Num
for key,value in dfctest[4].items():
    dfoutput['Critical Value (%s)'%key] = value
print(dfoutput)
```


Results of ADF Test:

Test Statistic	-1.633727e+01
P-value	3.020024e-29
#Lags Used	3.000000e+00
Number of Observations Used	9.360000e+02
Critical Value (1%)	-3.437356e+00
Critical Value (5%)	-2.864633e+00
Critical Value (10%)	-2.568417e+00


dtype: float64

Here the ADF test's p-value of 3.020962e-29 proves that the data is it is very unlikely that the data has a unir root so the data is stationary

```
In [66]: #
plt.figure()
sns.histplot(data_merge['WTISPLC'].dropna().pct_change(), kde=True)
plt.title(f"Distribution of Oil returns")
plt.show()
```



```
In [65]: fig = sm.qqplot(data_merge['WTISPLC'].dropna().pct_change(), line='q')  
plt.title('QQ-plot of Oil Returns')  
plt.show()
```



Lets do some exploration about whether the Oil returns data is normally distributed and we are doing this by a QQ-plot that shows that the data has big tails and this suggest that overall the data may look like it is normally distributed by the bell shaped histogram plot. Actually it is not because of the fat tails. For better prove of this hypothesis we will do the statistical test for normality check called Shapiro-Wilk test for normality. This will undeniably prove or deny the hipothesis mentioned above

```
In [48]: # Shapiro-Wilk Test
statistic, p_value = stats.shapiro(data_merge['WTISPLC'].dropna())
print('Shapiro-Wilk Test:')
print('Statistic:', statistic)
print('P-value:', p_value)
```

Shapiro-Wilk Test:
Statistic: 0.8222251534461975
P-value: 4.956003613292808e-31

Yet againd a p-value of 4.912296382340533e-31 for Shapiro-Wilk normality test proves the the data is not normally distributed. With such a small p-value we cannot accept the null hypothesis (which suggests that the data is normally distributed) so in this case we fail to accept it and we can say that we have statistically proved that the data is NOT normally distributed. Something else is very interesting and it is the statistic= 0.822 for similarity to a normal distribution. This suggest that the data is close to a normal distributio (the bell shape) but the fat tails fails this thesis

```
In [49]: from statsmodels.stats.diagnostic import acorr_ljungbox

# Perform Ljung-Box test up to lag 10
result = acorr_ljungbox(data_merge['WTISPLC'].dropna(), lags=10)


print(result)
```

	lb_stat	lb_pvalue
1	928.818987	5.333691e-204
2	1834.309316	0.000000e+00
3	2713.712826	0.000000e+00
4	3568.678890	0.000000e+00
5	4401.499486	0.000000e+00
6	5214.266084	0.000000e+00
7	6010.910499	0.000000e+00
8	6792.604186	0.000000e+00
9	7559.300054	0.000000e+00
10	8311.736911	0.000000e+00

With Ljung-Box Test we check for autocorrelation and prove that there is such with lag of 1.
(The lag 1 p-value of Ljung-Box Test is: 6.392185e-204 which is a lot lower than 0.05). This is also visible in the PACF plot but it is good to have it statistically proven

```
In [50]: data = data_merge['WTISPLC'].dropna()

# Calculate and plot rolling standard deviation (e.g., 20-period window)
rolling_std = data.rolling(window=20).std()
plt.plot(rolling_std)
plt.title('Rolling Standard Deviation of Oil Prices')
plt.show()
```



It is quite normal for Oil prices to have big spikes in volatility (expecially in war time periods) but also a big trough are quite normal at time of crisis where the demand and usage of oil is much lower than normally like in the 2008 Financial crisis and the Covid19 as well. So with the plot of the Rolling Standard Deviation of Oil Prices we can see exactly that: volatility clustering around periods of boom and bust (economy progress and crisis)

In [50]: